

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>.



deeplearning.ai

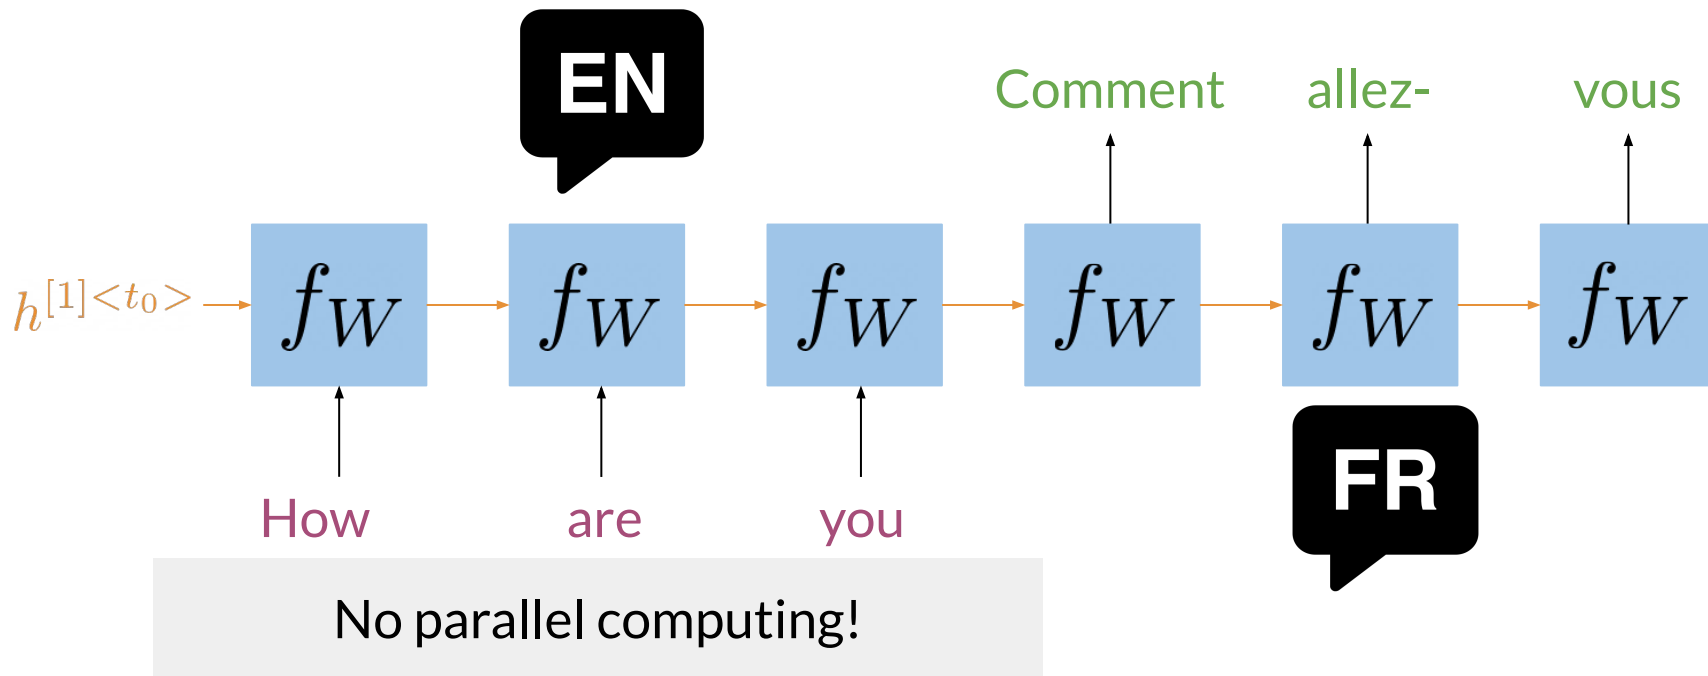
# Transformers vs RNNs

# Outline

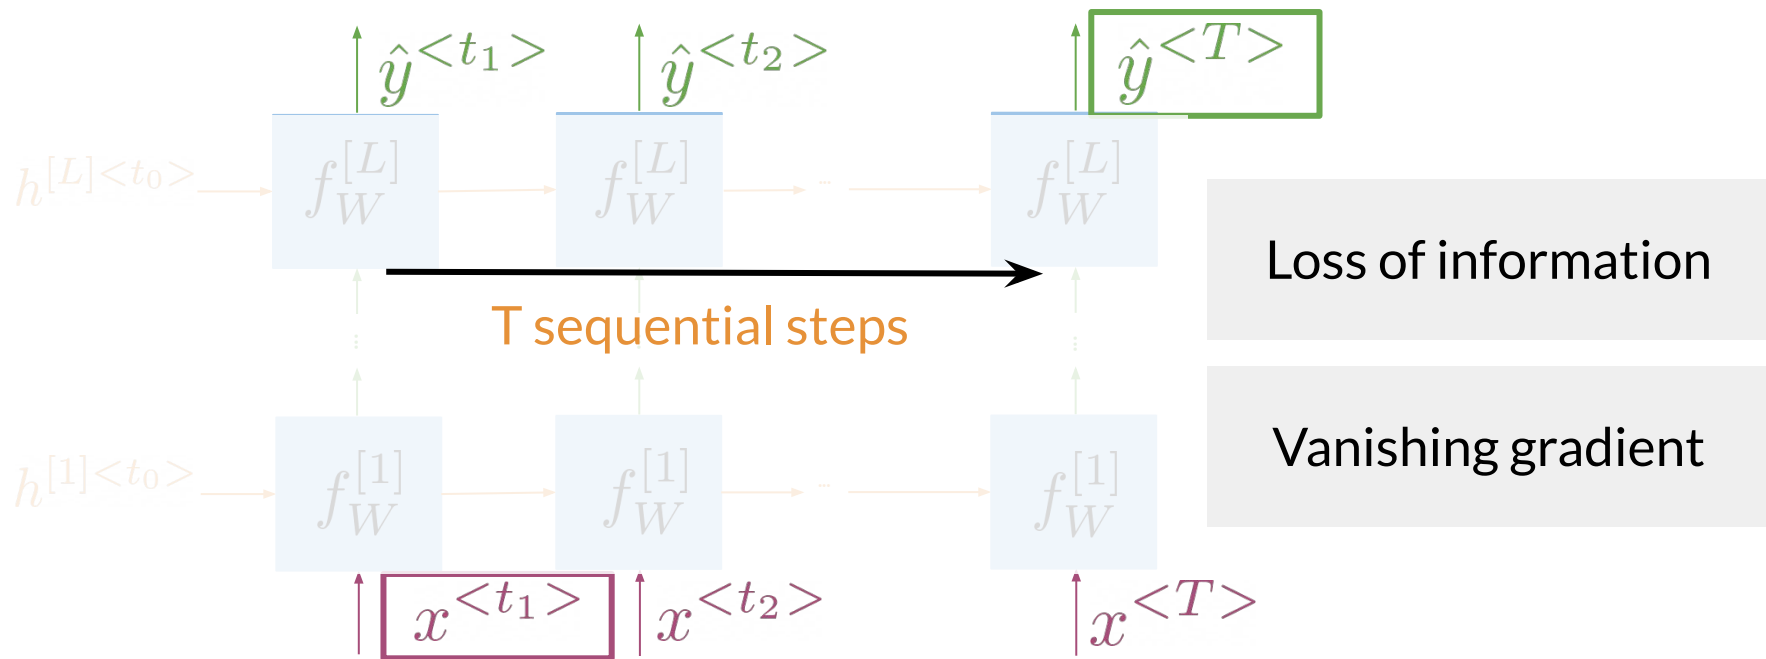
- Issues with RNNs
- Comparison with Transformers



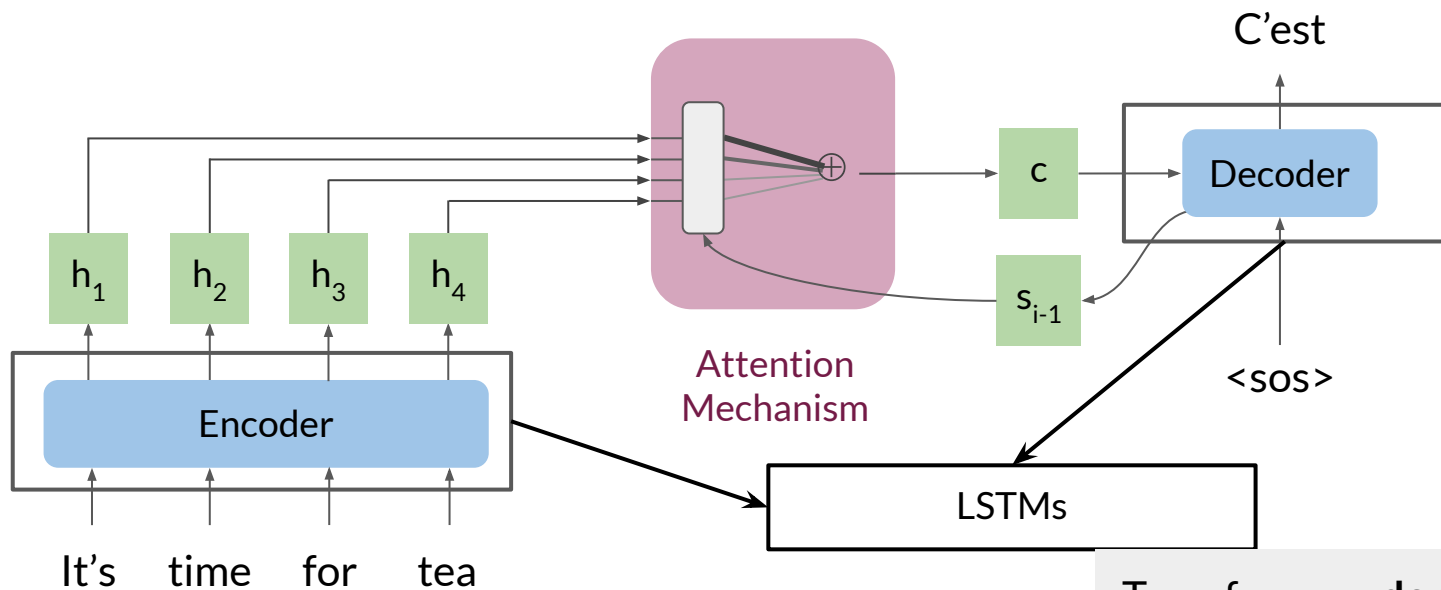
# Neural Machine Translation



# Seq2Seq Architectures



# RNNs vs Transformer: Encoder-Decoder



Transformers **don't** use RNNs, such as LSTMs or GRUs



deeplearning.ai

# Transformers Overview

# The Transformer Model

---

## Attention Is All You Need

---

**Ashish Vaswani\***  
Google Brain  
avaswani@google.com

**Noam Shazeer\***  
Google Brain  
noam@google.com

**Niki Parmar\***  
Google Research  
nikip@google.com

**Jakob Uszkoreit\***  
Google Research  
usz@google.com

**Llion Jones\***  
Google Research  
llion@google.com

**Aidan N. Gomez\* †**  
University of Toronto  
aidan@cs.toronto.edu

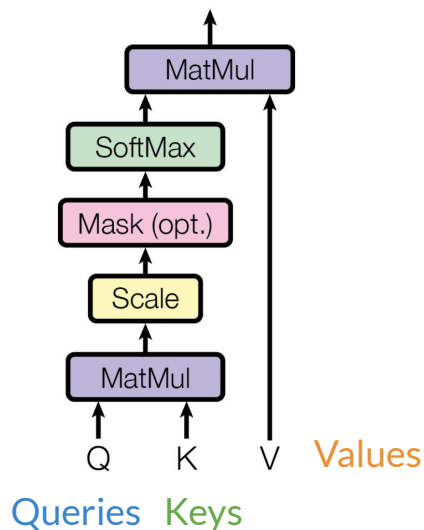
**Lukasz Kaiser\***  
Google Brain  
lukaszkaizer@google.com

**Illia Polosukhin\* ‡**  
illia.polosukhin@gmail.com

<https://arxiv.org/abs/1706.03762>



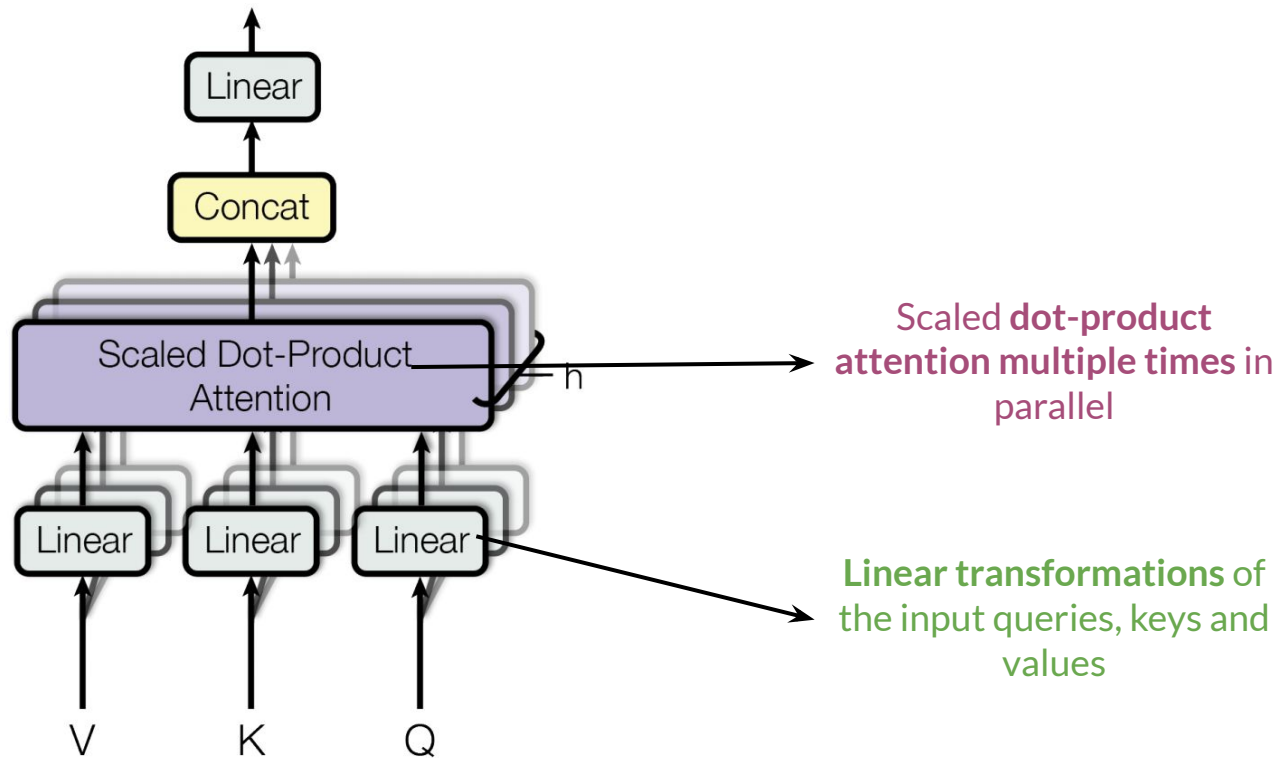
# Scaled Dot-Product Attention



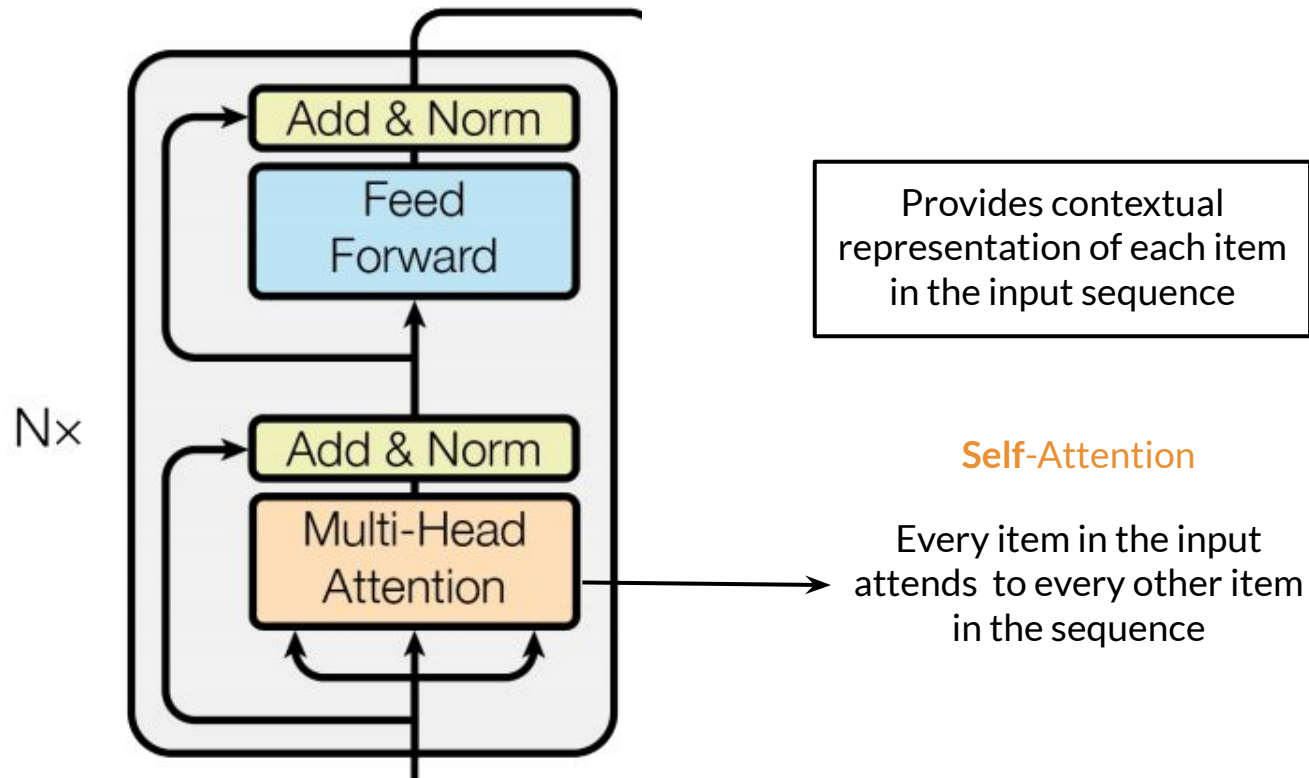
$$\text{softmax} \left( \frac{QK^{\top}}{\sqrt{d_k}} \right) V$$

(Vaswani et al., 2017)

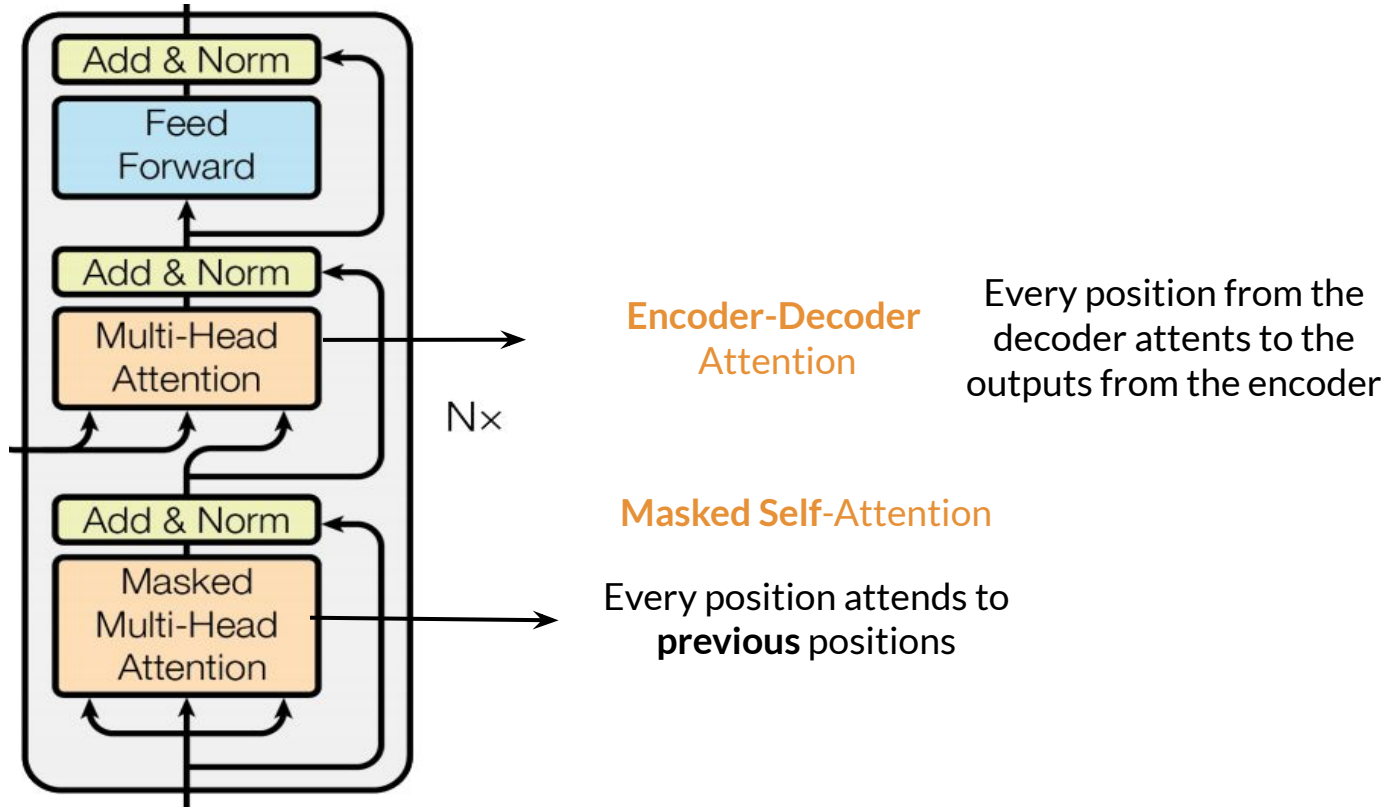
# Multi-Head Attention



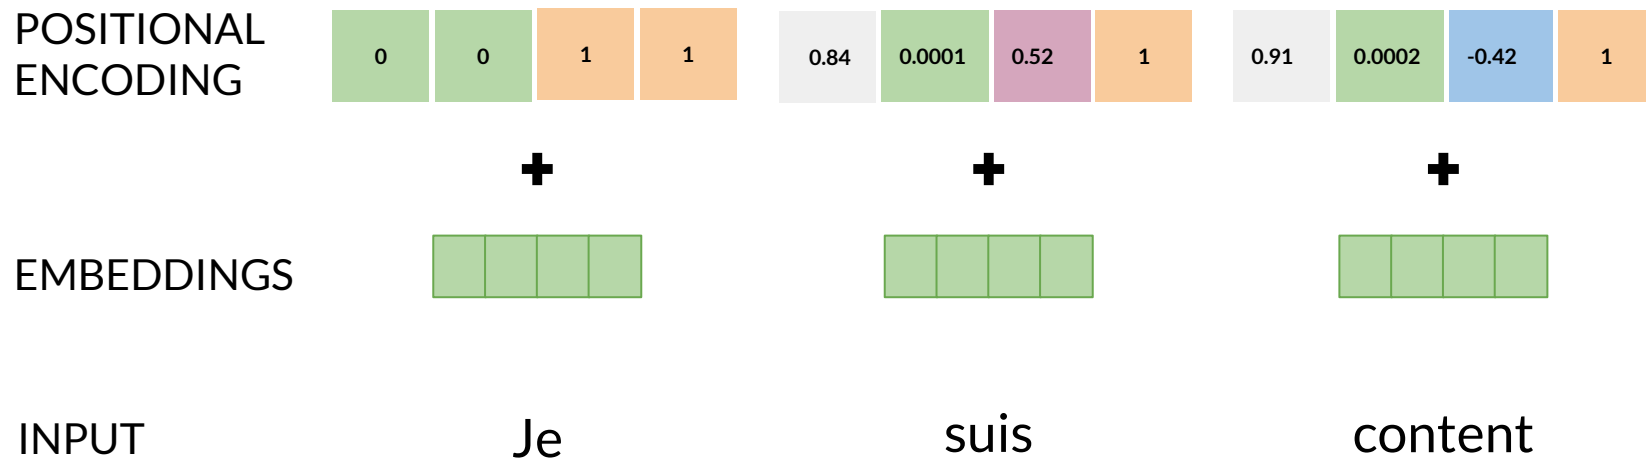
# The Encoder



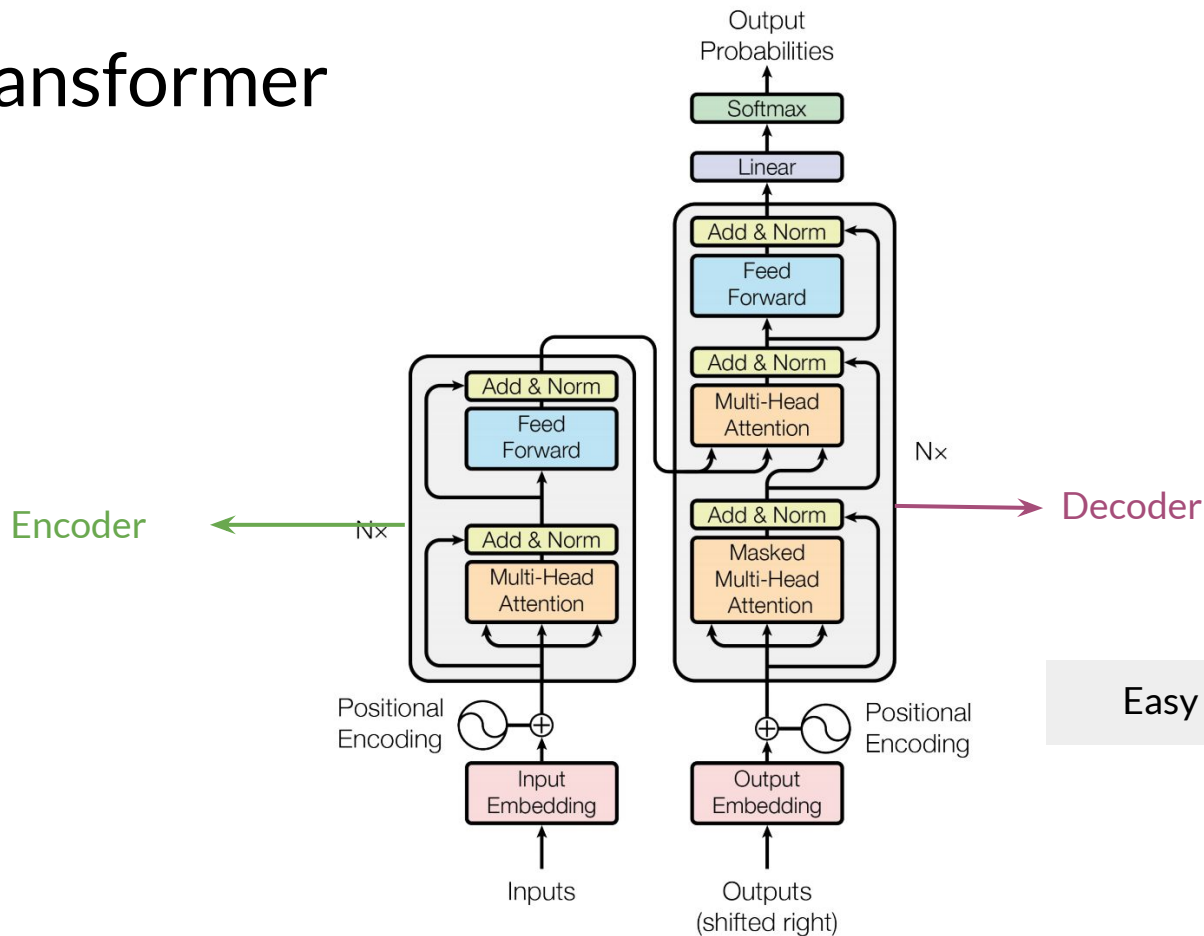
# The Decoder



# RNNs vs Transformer: Positional Encoding

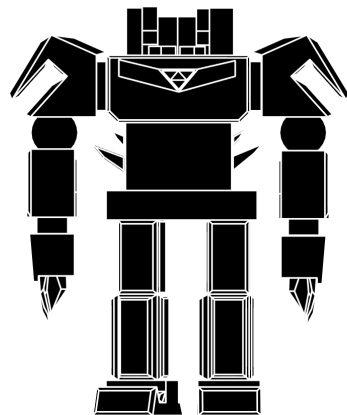


# The Transformer



# Summary

- In RNNs parallel computing is difficult to implement
- For long sequences in RNNs there is loss of information
- In RNNs there is the problem of vanishing gradient
- Transformers help with all of the above





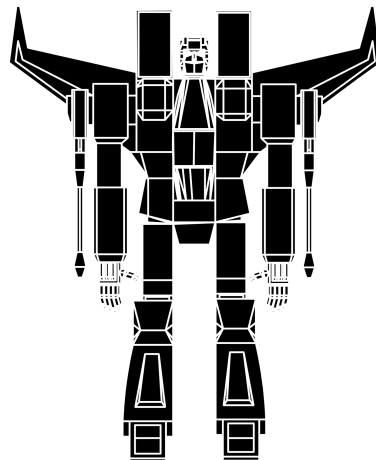
deeplearning.ai

# Transformer Applications



# Outline

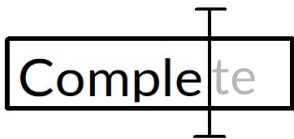
- Transformers applications in NLP
- Some Transformers
- Introduction to T5



# Transformer NLP applications



Text  
summarization



Auto-Complete

The	wind	blows	hard
Article	Noun	Verb	Adjective

Named entity  
recognition (NER)



Question  
answering (Q&A)

Translation



Chat-bots



Other NLP tasks

Sentiment Analysis  
Market Intelligence  
Text Classification  
Character Recognition  
Spell Checking

# State of the Art Transformers

Radford, A., et al. (2018)  
Open AI

Devlin, J., et al. (2018)  
Google AI Language

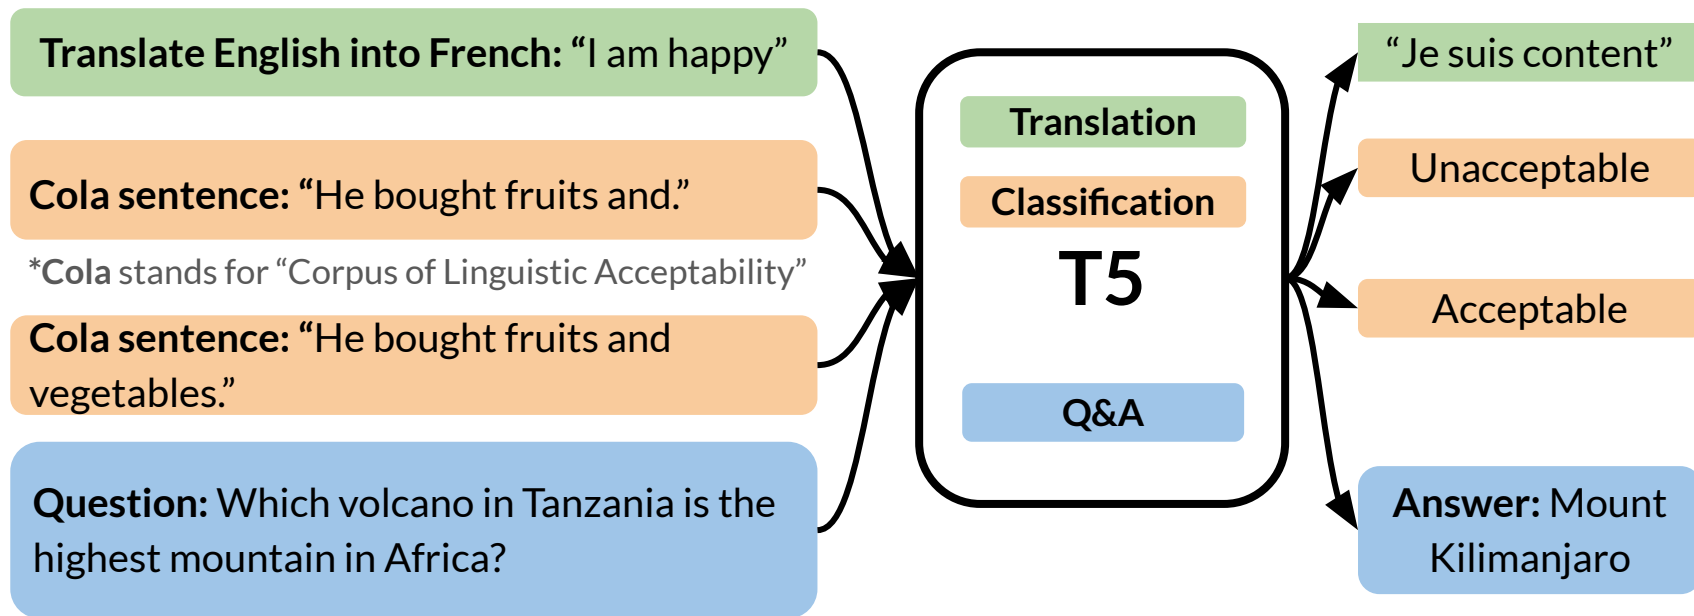
Colin, R., et al. (2019)  
Google

**GPT-2**: Generative Pre-training for  
Transformer

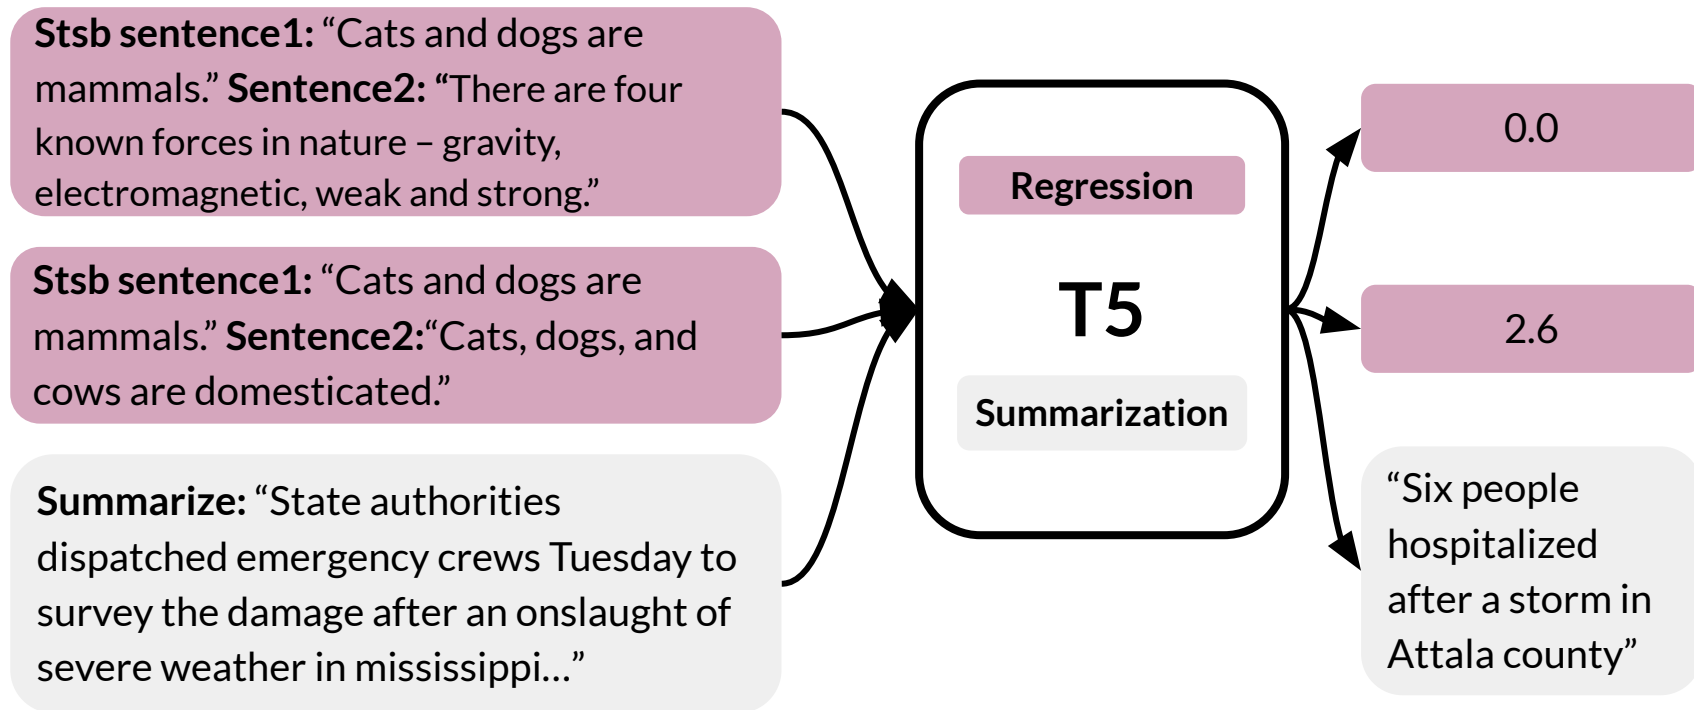
**BERT**: Bidirectional Encoder  
Representations from Transformers

**T5**: Text-to-text transfer transformer

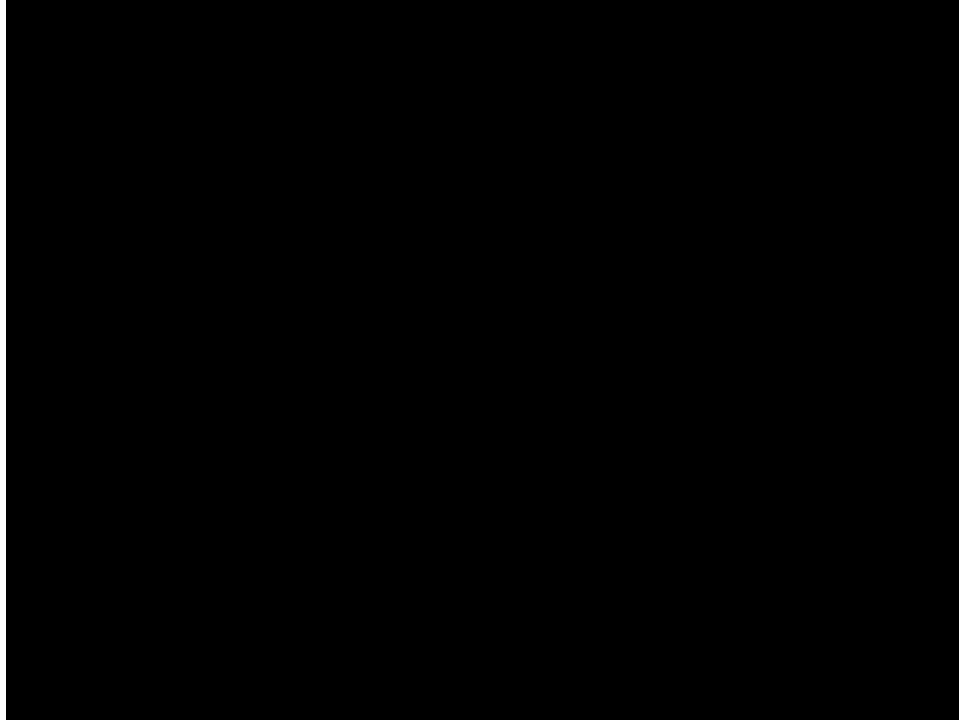
# T5: Text-To-Text Transfer Transformer



# T5: Text-To-Text Transfer Transformer

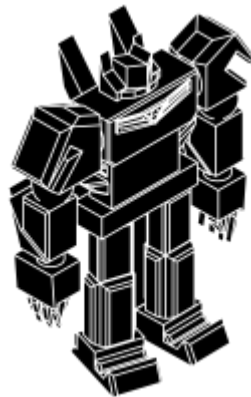


# T5: Demo



# Summary

- Transformers are suitable for a wide range of NLP applications
- Some transformers include GPT, BERT and T5
- T5 is a powerful multi-task transformer





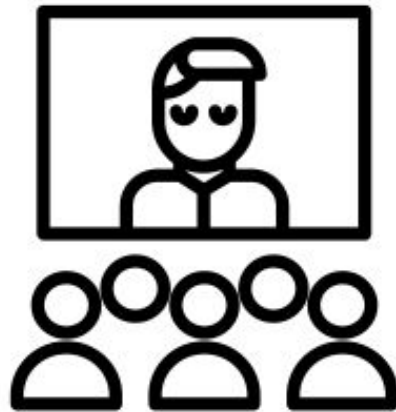
deeplearning.ai

# Scaled Dot-Product Attention

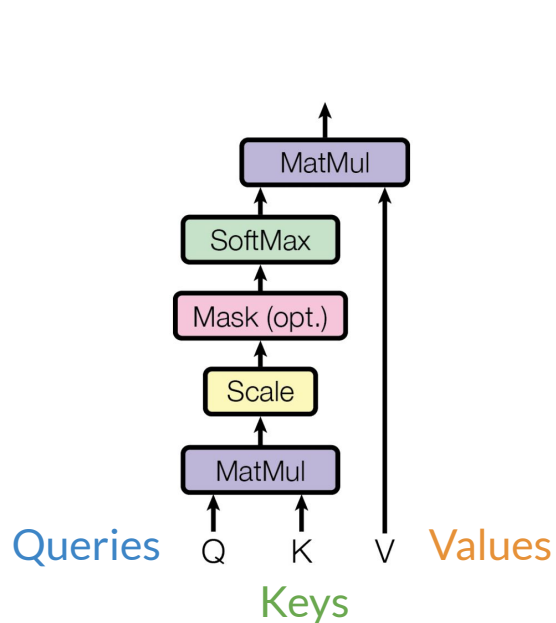


# Outline

- Revisit scaled dot product attention
- Mathematics behind Attention



# Scaled dot-product attention



(Vaswani et al., 2017)

Weights add up to 1

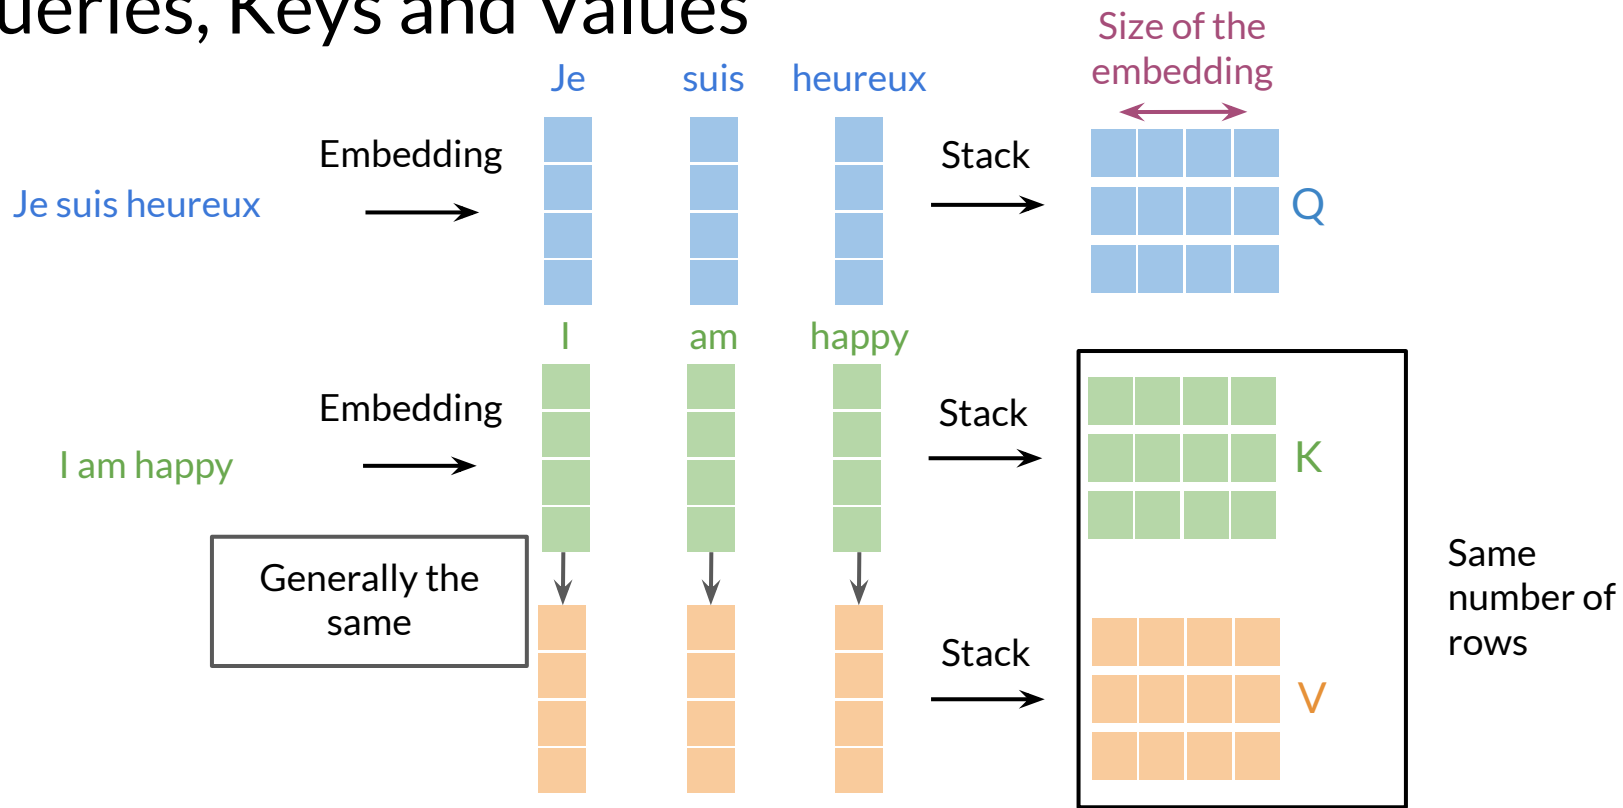
Improves performance

$$\text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

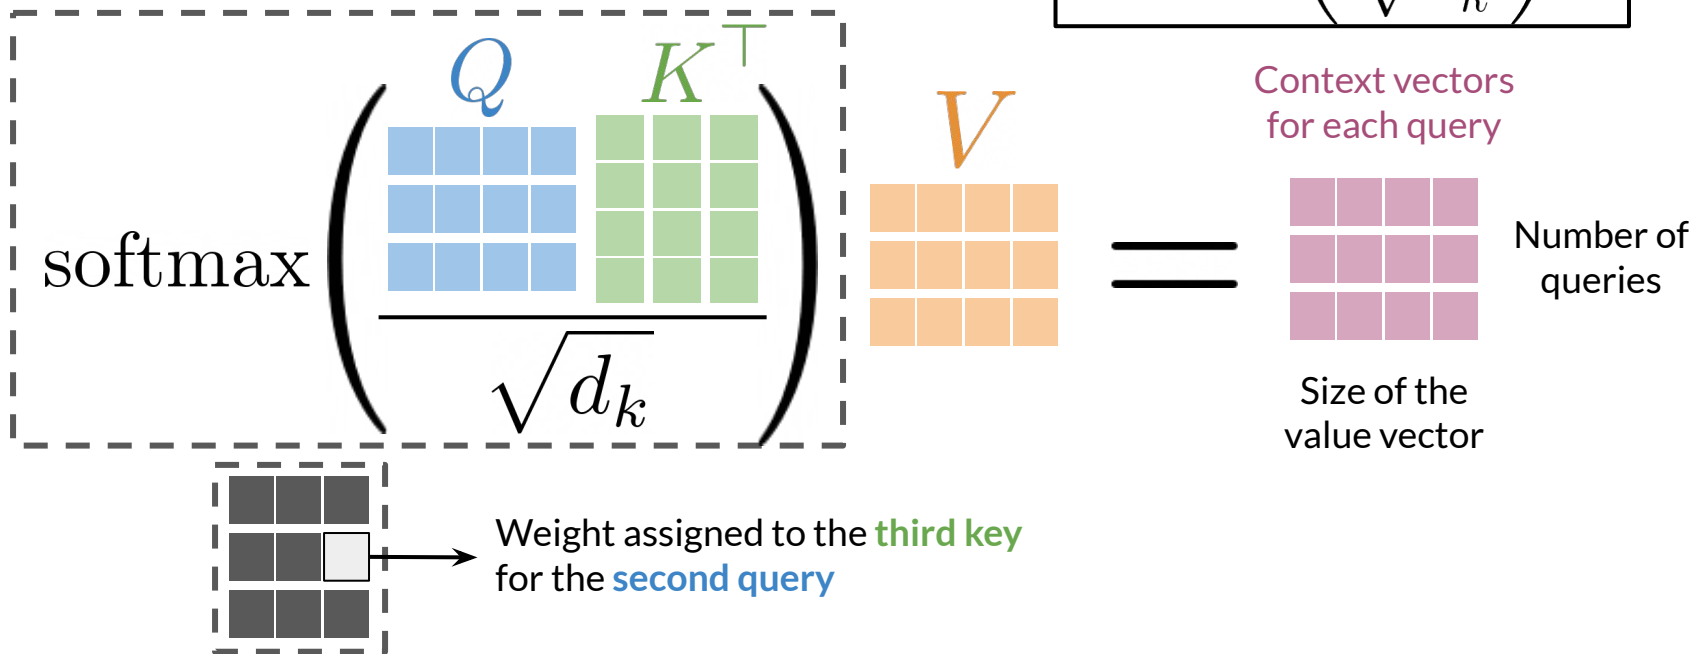
Weighted sum of values  $V$

Just two matrix multiplications  
and a Softmax!

# Queries, Keys and Values



# Attention Math



# Summary

- Scaled Dot-product Attention is essential for Transformer
- The input to Attention are queries, keys, and values
- GPUs and TPUs





deeplearning.ai

# Masked Self-Attention

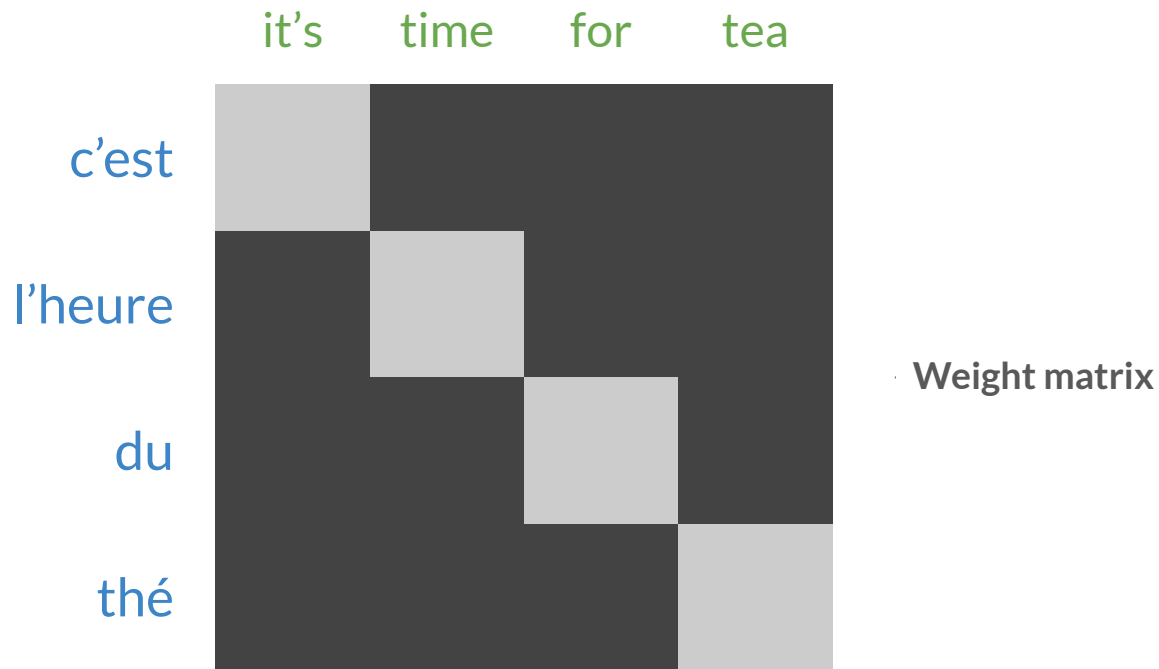
# Outline

- Ways of Attention
- Overview of masked Self-Attention



# Encoder-Decoder Attention

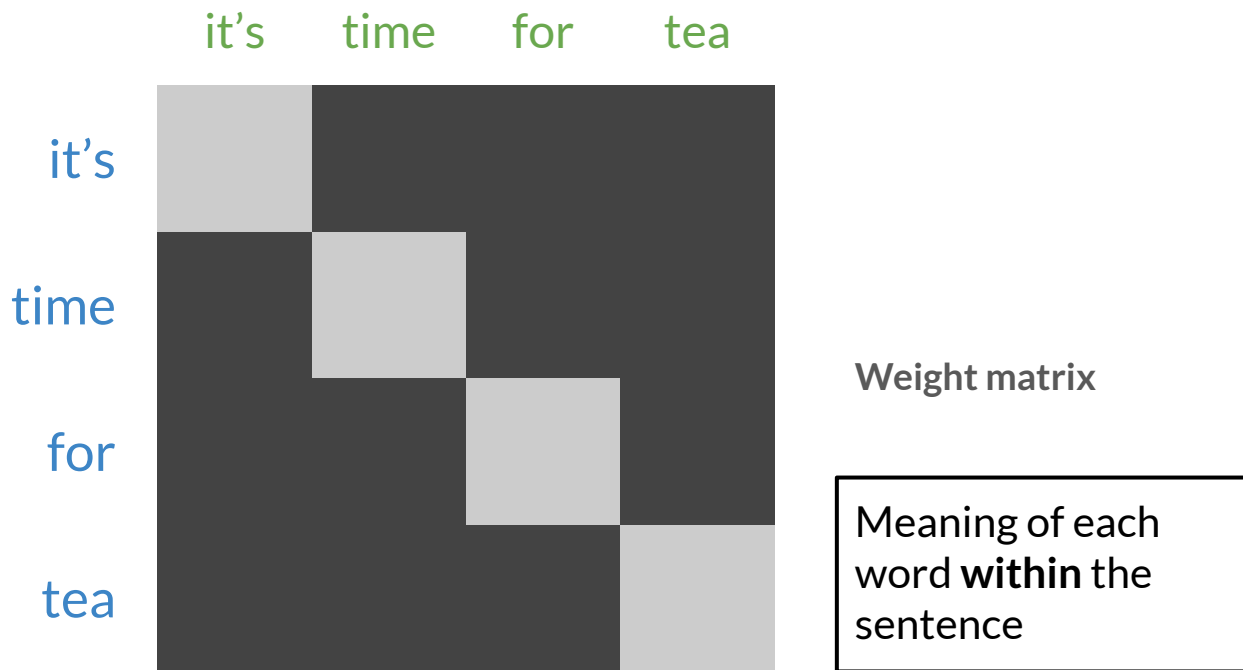
Queries from one sentence, **keys** and **values** from another





# Self-Attention

Queries, keys and values come from the **same sentence**



# Masked Self-Attention

Queries, keys and values come from the **same sentence**. Queries don't attend to future positions.



# Masked self-attention math

Diagram illustrating the masked self-attention mechanism:

The formula for masked self-attention is:

$$\text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} + \text{Mask} \right) V$$

Where:

- $Q$  (Query) is a 3x4 matrix of blue squares.
- $K^T$  (Key Transpose) is a 4x3 matrix of green squares.
- $V$  (Value) is a 3x4 matrix of orange squares.
- $\sqrt{d_k}$  is the square root of the key dimension.
- The **Mask** is a 3x3 matrix where non-zero values (0) represent valid attention weights and pink squares represent "Minus infinity".

Mask Matrix (3x3):

0	Minus infinity	Minus infinity
0	0	Minus infinity
0	0	0

Weights assigned to future positions are equal to 0

Future Position Mask Matrix (3x3):

Dark Gray	0	0
Dark Gray	Dark Gray	0
Dark Gray	Dark Gray	Dark Gray

Legend: Pink square  $\rightarrow$  Minus infinity

# Summary

- There are three main ways of Attention: Encoder/Decoder, self-attention and masked self-attention.
- In self-attention, queries and keys come from the same sentence
- In masked self-attention queries cannot attend to the future





deeplearning.ai

# Multi-head Attention

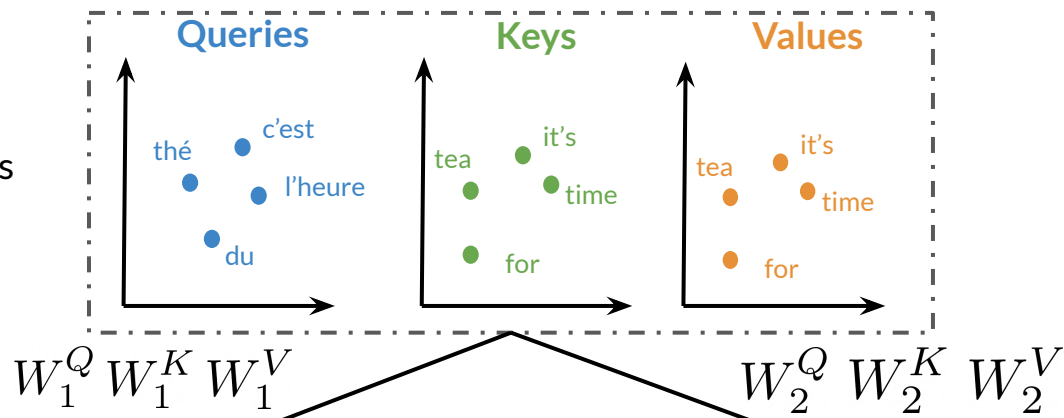
# Outline

- Intuition Multi-Head Attention
- Math of Multi-Head Attention

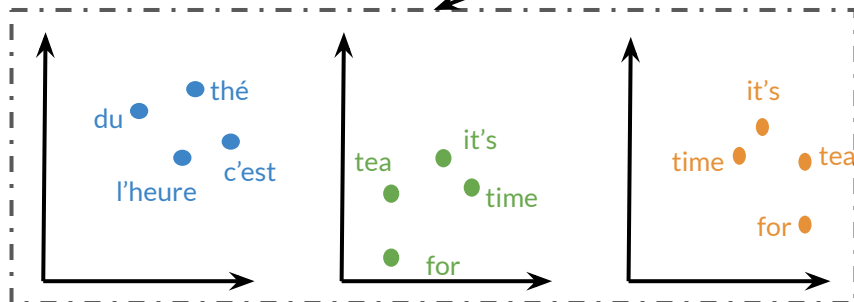


# Multi-Head Attention - Overview

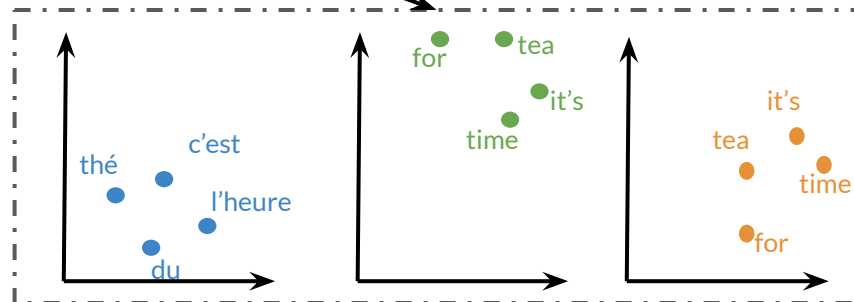
Original Embeddings



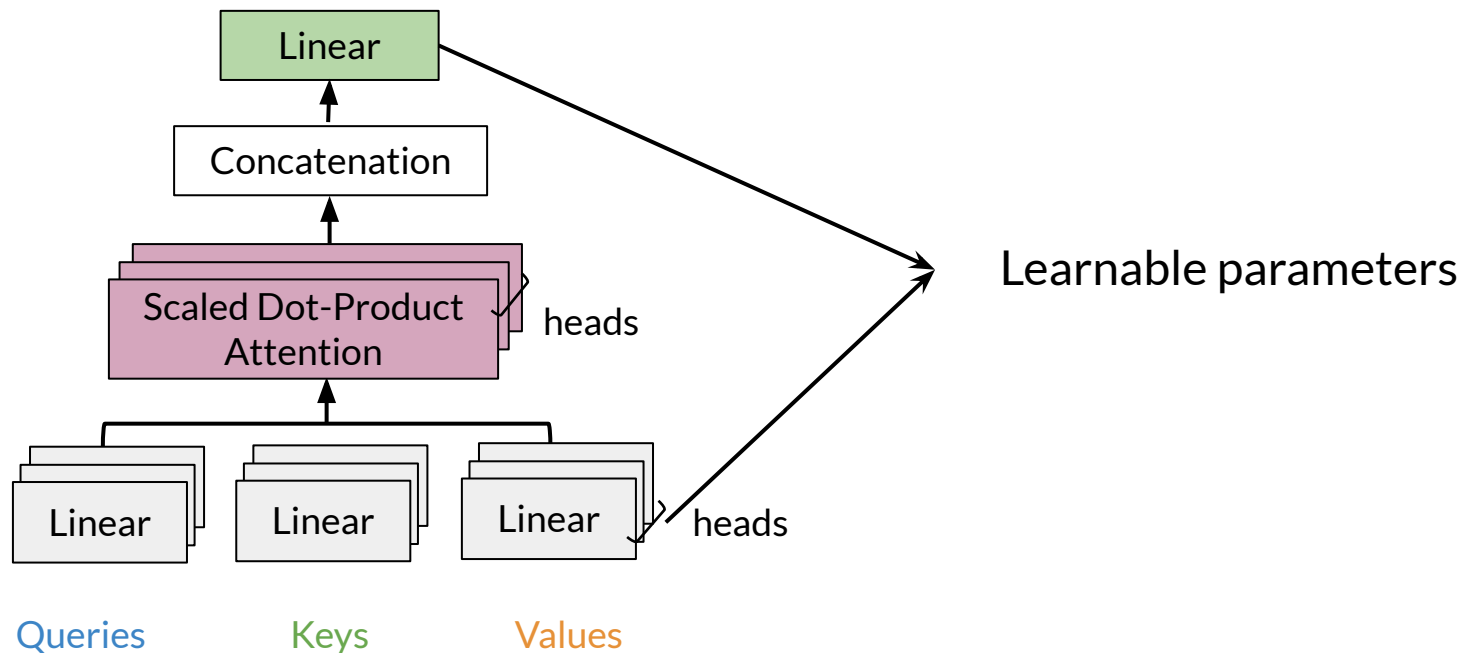
Head 1



Head 2



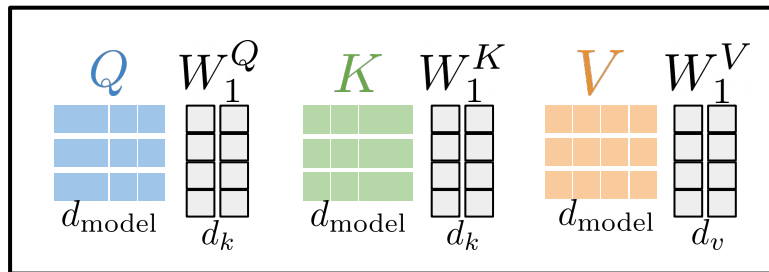
# Multi-Head Attention - Overview





# Multi-Head Attention

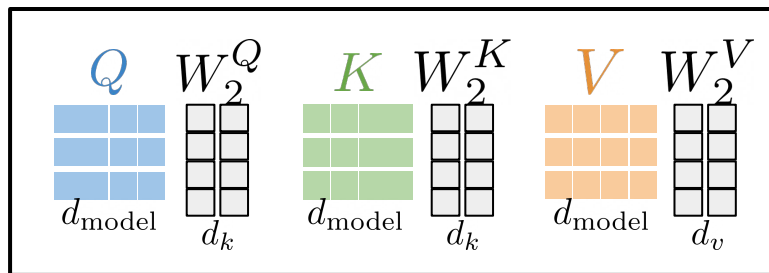
Head 1



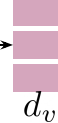
Attention



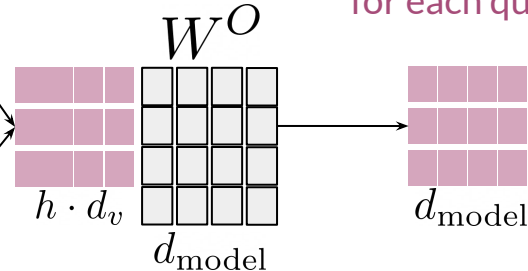
Head 2



Attention



Concat



Context vectors  
for each query

$d_{\text{model}}$ : Embedding size

Usual choice of dimensions  
 $d_k = d_v = d_{\text{model}}/h$

# Summary

- Multi-Headed models attend to information from different representations
- Parallel computations
- Similar computational cost to single-head attention





deeplearning.ai

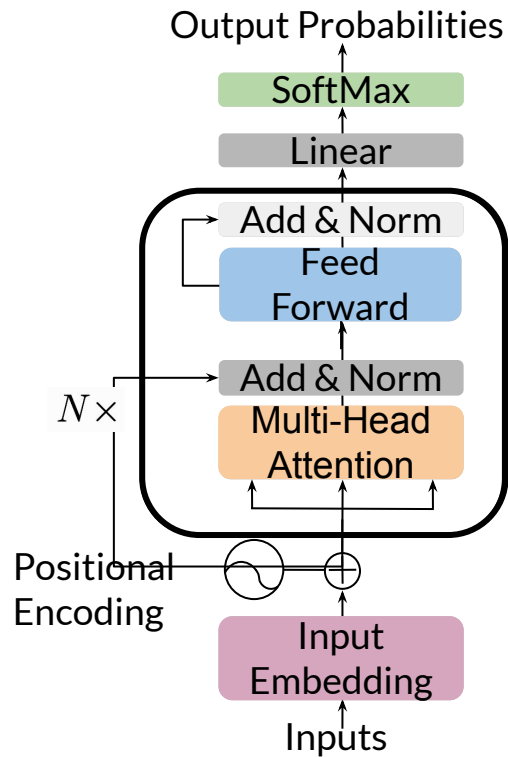
# Transformer decoder

# Outline

- Overview of Transformer decoder
- Implementation (decoder and feed-forward block)



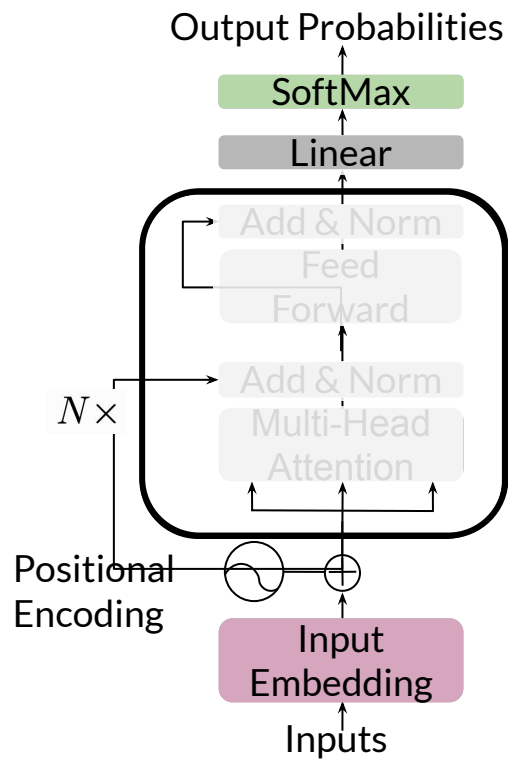
# Transformer decoder



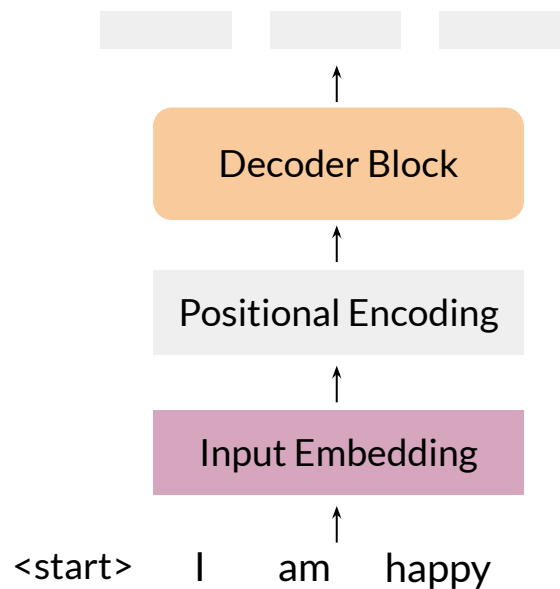
## Overview

- input: sentence or paragraph
  - we predict the next word
- sentence gets embedded, add positional encoding
  - (vectors representing  $\{0, 1, 2, \dots, K\}$ )
- multi-head attention looks at previous words
- feed-forward layer with ReLU
  - that's where most parameters are!
- residual connection with layer normalization
- repeat N times
- dense layer and softmax for output

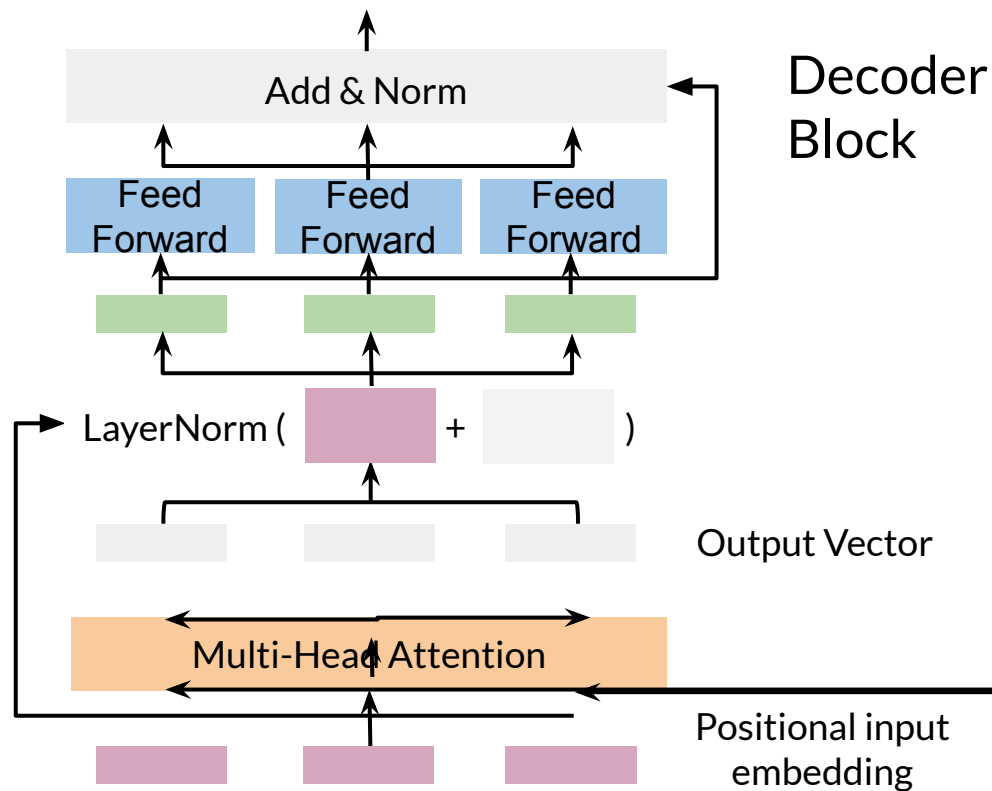
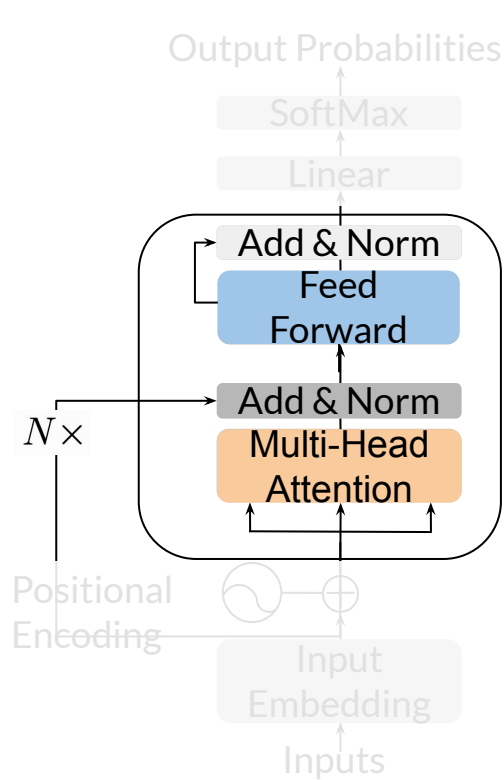
# Transformer decoder

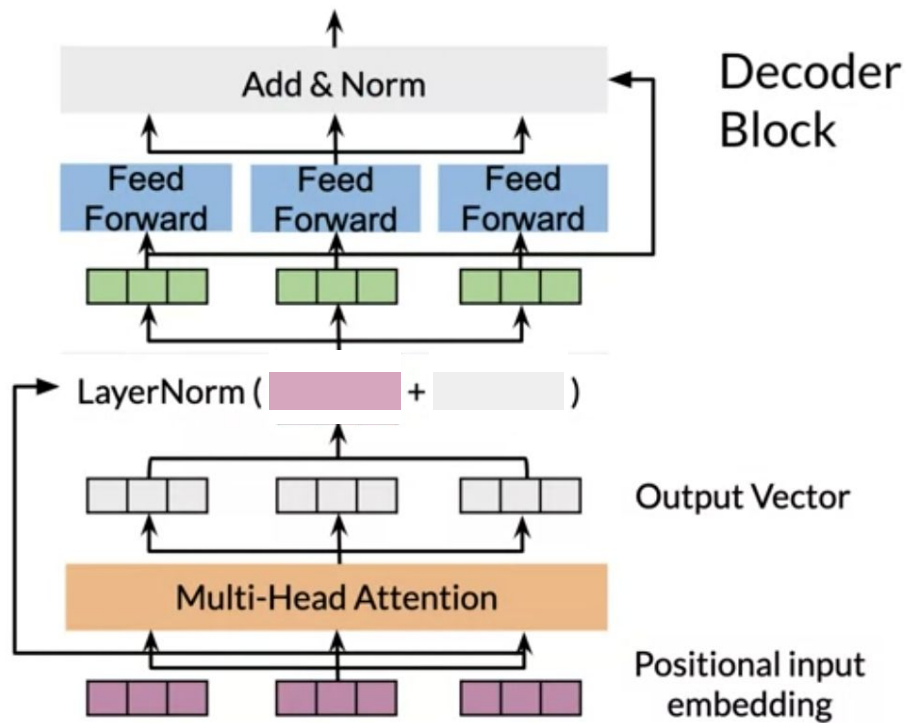


## Explanation



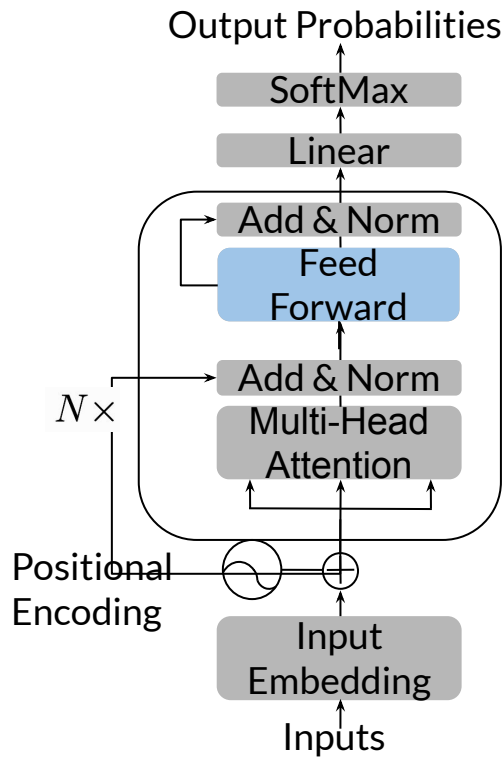
# The Transformer decoder



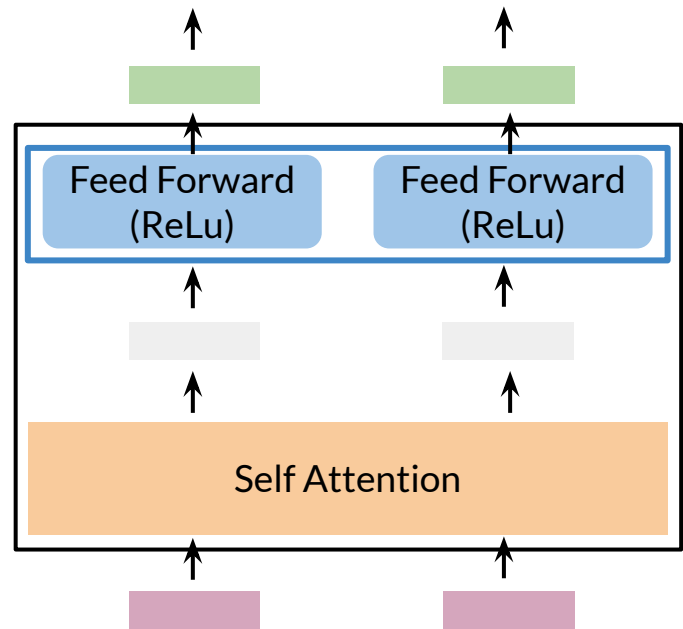




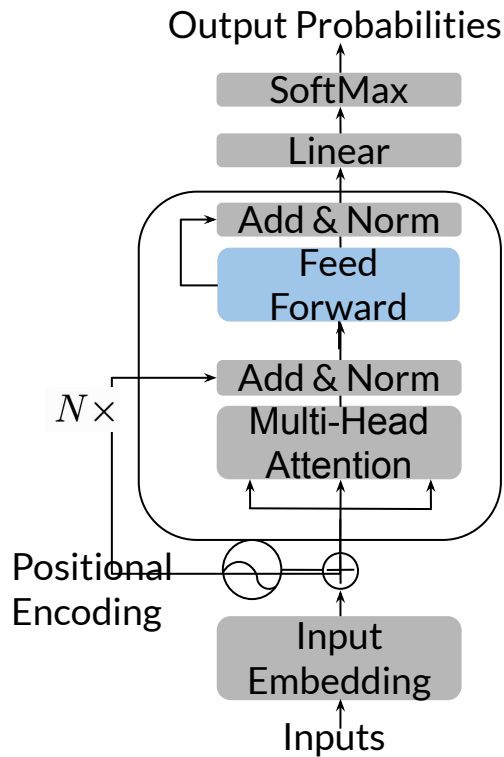
# The Transformer decoder



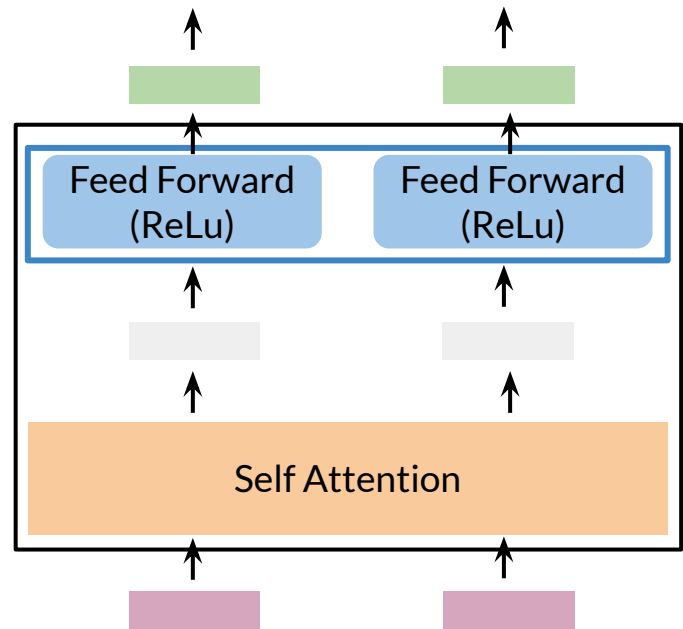
## Feed forward layer



# The Transformer decoder



## Feed forward layer



# Summary

- Transformer decoder mainly consists of three layers
- Decoder and feed-forward blocks are the core of this model code
- It also includes a module to calculate the cross-entropy loss



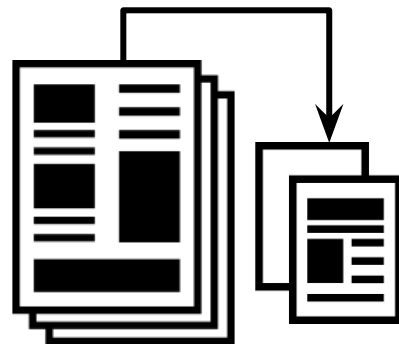


deeplearning.ai

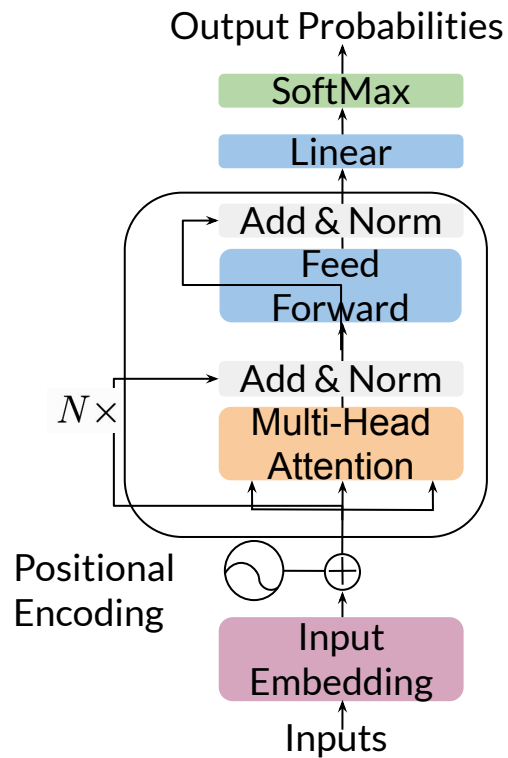
# Transformer summarizer

# Outline

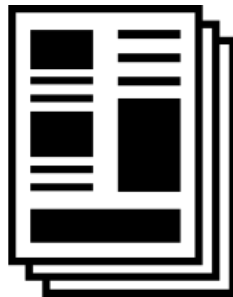
- Overview of Transformer summarizer
- Technical details for data processing
- Inference with a Language Model



# Transformer for summarization



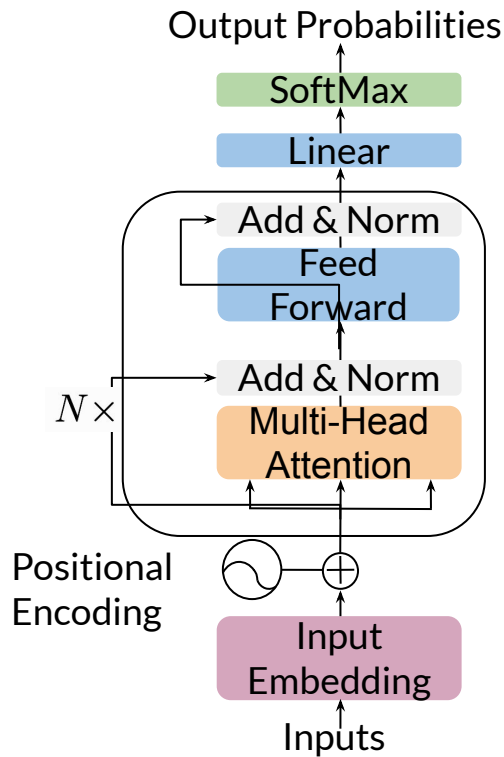
Input



Output:  
Summary



# Technical details for data processing



## Model Input:

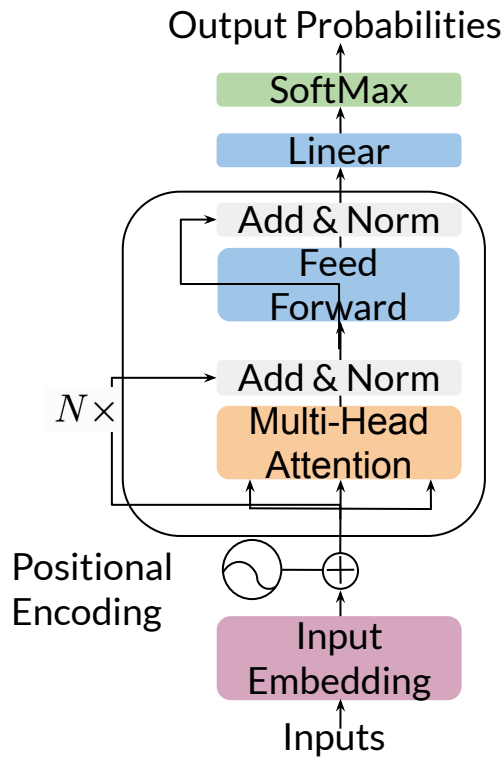
ARTICLE TEXT <EOS> SUMMARY <EOS> <pad> ...

## Tokenized version:

[2,3,5,2,1,3,4,7,8,2,5,1,2,3,6,2,1,0,0]

Loss weights: 0s until the first <EOS> and then 1 on the start of the summary.

# Cost function



## Cross entropy loss

$$J = -\frac{1}{m} \sum_j^m \sum_i^K y_j^i \log \hat{y}_j^i$$

$j$  : over summary

$i$  : batch elements





# Inference with a Language Model

## Model input:

[Article] <EOS> [Summary] <EOS>

## Inference:

- Provide: [Article] <EOS>
- Generate summary word-by-word
  - until the final <EOS>
- Pick the next word by random sampling
  - each time you get a different summary!

# Summary

- For summarization, a weighted loss function is optimized
- Transformer Decoder summarizes predicting the next word using
- The transformer uses tokenized versions of the input

